

Your Agent Found 5 Bugs. 4 Were Imaginary.

How to measure trust in an AI testing agent — the results quadrant



1 real bug. 4 artifacts.

Our DevAssure O2 run found a real bug — whitespace-only posts could be submitted. Then it produced 4 more findings, all artifacts of its own actions: it injected data, then "discovered" the injection.

1/5

20% confirmation

80% false-discovery

The 4 cells of agent results

AGENT REPORTED

TP

real bug, correctly reported

AGENT REPORTED

FP

hallucinated bug

AGENT SILENT

FN

missed a real bug

AGENT SILENT

TN

correctly stayed quiet

Demos show TP. The other 3 cells are what demos hide.

FP and FN are the quiet cells you need to measure

4 failure modes, 4 probes

The hallucination · FP

reports a bug that isn't there

 review gate

The miss · FN

silent on a real bug

 fault injection · 34/34

The healed test · FN

green, but checks a string not behavior

 diff every change

The blind first-match · FN

picks the first target, ambiguity unseen

 decoy targets

Each failure mode needs its own probe

RED BUILD

Forces a decision

clear signal

GREEN REPORT

From a black box

"no risk detected"

**A red build forces a decision.
A green report from a black box does not.**

Test the tool that tests for you — or pair it with an oracle you trust.

Full breakdown — article in comments

Two experiments (DevAssure O2 + a commercial AI-QA platform decoy),
the **4 failure modes with probes**, and a **Monday checklist**.



**Do you track confirmation rate for agent-reported bugs — or just
count how many it found?**

#AIAgents #GenAITesting #MutationTesting