



QUALITY OPERATING MODEL

How to Evaluate Any **AI-QA Vendor** in 5 Scenarios

Break the testing tool on purpose.

THE PROBLEM

A tool that fails red costs you an hour. A tool that **lies green costs you the release.**

5/5 green at 100% confidence, 0% risk — **with two identical login forms on the page.** The tool saw the diff. It never flagged the ambiguity.

THE METHOD — 5 SCENARIOS, ONE QUESTION PER RUN

Mutation matrix — 5 scenarios


Same login flow, 5 scenarios, one question per run: what does the tool report?

	1	2	3	4	5
	Baseline clean page confirm green	Locator drift id / class changed same semantics	Weak decoy duplicate text on page	Strong decoy identical form same style	Product regression label changed locator stable
EXPECTED report	GREEN	GREEN / heal	FLAG	FLAG	CATCH
ACTUAL AI report	GREEN	GREEN	GREEN	GREEN	GREEN

When every row is green, green means nothing. **Silent green = silent false negative.**

When every row is green, **green means nothing.**


TWO VENDORS, TWO FAILURE SHAPES



Silent green vs honest death. One hides the defect. The other dies loudly — and tells you why.

THE LOOP CLOSED

Breaking it shipped a feature



Reported the gap instead of writing the tool off. Now a passed run flags its own warnings. **Green stayed green — it stopped being silent.**

THE PLAYBOOK

Evaluate any vendor **in 4 steps**

1. Import a small suite you fully understand — 5 tests beat 500.
2. Run the 5 scenarios; demand per-scenario outputs.
3. Red flags: silent green, heal-without-review, confidence without methodology.
4. Report gaps first. Tools that ship fixes in days are keepers.

YOUR TURN

When did you last **break your testing tool** on purpose?

Full method + numbers in the article. Part of the **Quality Operating Model** series.

Victor Ematin · AI Quality Engineering Lead